

# CmpE 275

---

## Section: Concepts of messaging

Sequential, Spatial, and Temporal Decoupling – meh...right?

# On track

- Agenda (baseline of all distributed computing)
  - ♦ Defining distributed computing
  - ♦ On messaging
    - Synchronous network systems (messaging)
    - Asynchronous systems
    - Example: Coordination through Leader Election Algos
- Key points
  - ♦ Synchronous and Async
  - ♦ Message design strategies
  - ♦ Spatial/Temporal decoupling

# Beyond the metal lab

- The basic mechanics of the socket lab introduces challenges that we will explore in depth. In the bare-metal we constructed a single client-server (two-tier). Consider:
  - ◆ Data representation
  - ◆ Long chains or cycles
  - ◆ Hot spots
  - ◆ Software – Abstraction vs. complexity
    - more features, more implementation hiding
  - ◆ Failure and recovery – lost connections, data
  - ◆ Velocity, Voracity, Volume – the three-Vs

# Complex data: Overloading

- How to represent data (encode)?
  - ◆ Attachments (e.g., Multi-part)
    - Data is stored as a separate document attached to the request (upload, download a file) or referenced URL (common storage)
  - ◆ Name-value pairs (attributes, parameters), tuples
    - Form-like (e.g, firstname-value, lastname-value)
    - Converting to name-value pairs is difficult if the data is complex and/or deeply nested – representing graphs of data (hierarchical), the data is flattened (row-column)
    - How to support inheritance? Data changes (releases)?

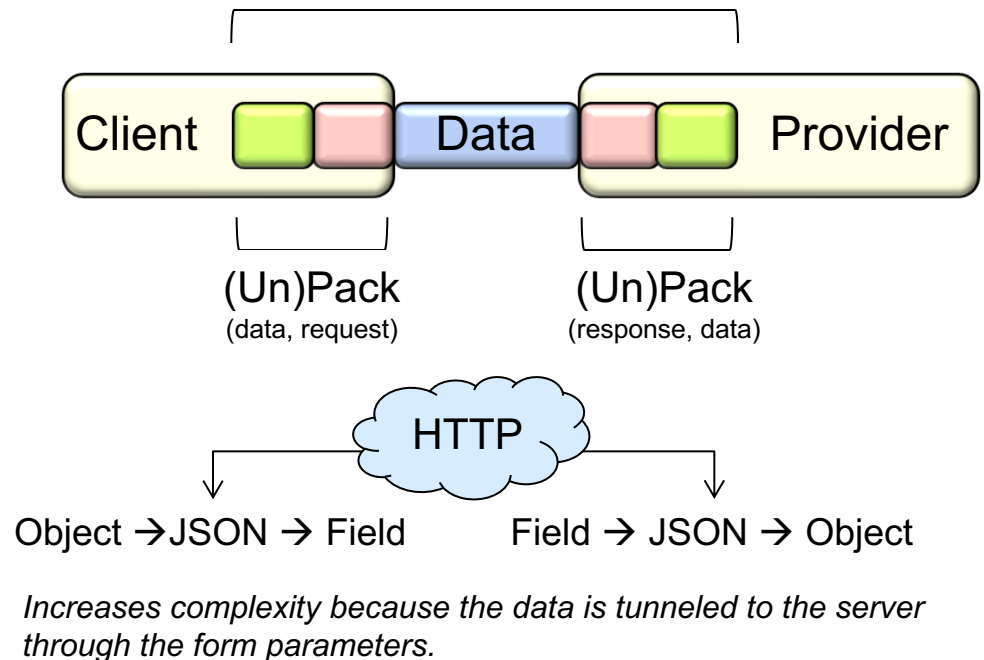
# Data payloads Representations

## (JSON, XML -> REST,...)

- Delivering complex data structures using overloading

- ◆ Type overloading
  - application/json, application/xml
  - Encoding data for the client and/or server to process
- ◆ Value overloading
  - Overloading POST or PUT

Netty/PB or  
REST/HTTP (XML, JSON) or  
SOAP/WSDL (XML, MTOM)



# Messaging

# Message-based communication

- The Message Queue (MQ) model
  - ♦ You have studied MQ as web services, JMS, ...
  - ♦ What functionality does it supply?  
What are the basic concepts?
- Can you decompose the MQ model into a sequence diagram?

# Is distributed the same as parallel as concurrent?

- Parallel
  - ♦ simultaneous, independent execution of tasks
- Concurrent
  - ♦ scheduled cooperative (interleaved) execution where (typically) only one thread is active.
- Distributed
  - ♦ Parallelization across processes in asynchronous designs



# What are the incentives to use distributed architectures?

1. Scaling
  - a. CPU
  - b. Memory
  - c. I/O
2. Failure-Recovery
3. Performance
4. Isolation



Image: Why? by Myles! on Flickr

# Sequence decoupling

- Parallel (Sequence decoupling)
  - ◆ Parallelization of a request into smaller components that can be processed concurrently.
  - ◆ Behavior is both synchronous and asynchronous
  - ◆ Advantages
    - Processing large amount of information that otherwise would be difficult or inefficient in a serial algorithm

# Spatial decoupling

- Location/Space (Spatial decoupling)
  - ◆ Interactions are not limited to the current process space, computer, system
  - ◆ The client and server are not required to co-exist on the same server, OS, and language
  - ◆ Advantages
    - Server-side scaling – architecture can change without affecting the producer

# Temporal decoupling


- Time (Temporal decoupling) – asynchronous behavior also frees processes to act
  - ◆ No timing dependencies between producer (client) and consumer (server). The consumer is not required to act immediately when a message is produced
  - ◆ Advantages
    - Consumer can defer processing a message. This allows the consumer to apply QoS and fair scheduling practices
    - Partially supports a partitioned network (why only partially?)

# Communication network architecture building blocks


- Building a network architecture is basically the combination of a few key concepts
  - ◆ Message passing patterns
  - ◆ Queuing behavior
  - ◆ Overlay network construction

# Types of message passing


- Message passing generally falls to the following patterns



pt-to-pt




Multicast




Broadcast

# Queue behavior: FIFO, LIFO, Selective


- enqueue and dequeue behavior of stacks or queues



FIFO




LIFO



Selective

# Patterns of network organization (overlay networks)

- Rings
- Hubs
- Bus
- General Graphs



How do these designs affect communication?


- Configured as
  - ◆ Centralized
  - ◆ Decentralized



# Asynchronous Messaging

# Synchronous message processing is similar in behavior between two classes


- From within a method (same process)
  - ♦ E.g., `A.do() { b = B.new(); b.something(); }`




if  $dt \ll \text{acceptable liveliness (AL)}$  then the solution is good. **What if  $dt \gg \text{AL}$ ?**

# Consider a ring network in synchronous communication

- Given 4 Nodes (A, B, C, D)
- Execution cycle for 'A' to send one message



1.sendTo2(m)




2.getMsg(m) {2.sendTo3(m);}



...

# Synchronous chained messaging does not have a constant dt ( $dt = \sum (dt_i + \text{Latency}_i)$ )



$$dt_{\text{Liveliness}} = dt_A + dt_B + dt_C + dt_d + (6 * \text{Latency})$$

---

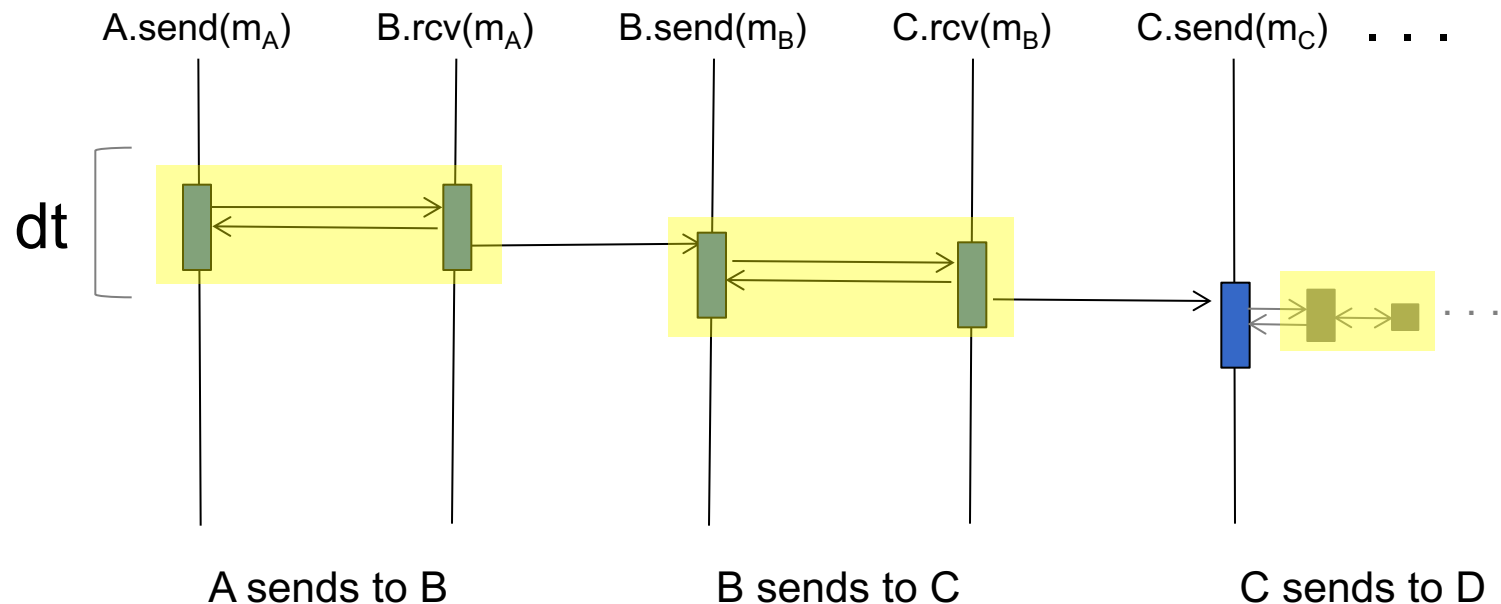
Technologies utilizing synchronous messaging include Java RMI, CORBA, Java EE 6 (EJB), Web Services, but not our socket toolkit!

# Asynchronous (Non-blocking) inter-process communication provides ~constant dt

$$(dt = dt_i + \text{Latency}_i)$$

- Liveliness from Node A's perspective

$$dt_{\text{Liveliness}} = dt_A + (1 * \text{Latency})$$



# Synchronous vs. Asynchronous problem solving

- Given
  - ♦ 5 nodes (processes) arranged in a ring
  - ♦ Each node is equidistant (1 m) from each other
  - ♦ A complete request uses all nodes (A-B-C-D)
  - ♦ A node can only process one request at a time
  - ♦ The latency between nodes is 250 msec
  - ♦ Nodes A and C are Intel XEON (4 cpu 8 core servers)
  - ♦ Nodes B and D are AMD (2 cpu 12 core servers)
  - ♦ Each node requires 1 min 30 sec to process a request
- For asynchronous and synchronous calculate
  1. What is  $dt$  for node A?
  1. How long does it take the network to process a request?

# Approaches associated with asynchronous messaging

- In the previous problem, a process could only operate on a single message at a time.
  - ♦ In a synchronous configuration, this was okay as blocking 'fit within' the design
- For asynchronous messaging (non-blocking requests), how do we provide single message processing while accepting multiple requests?


# Socket-per-thread

- If our server receives messages and for each message creates a thread
  - ♦ Like our socket-toolkit – right?
- Okay, so what is the concern with a thread per request?



# Socket to thread pool solution


- Limit the number of threads to minimize thread creation/execution



What about `m4` , `m5` , `mj` ?

# Queuing of inbound messaging provides a buffer between the receiving process (doing the work) and the sender (client requesting the work)

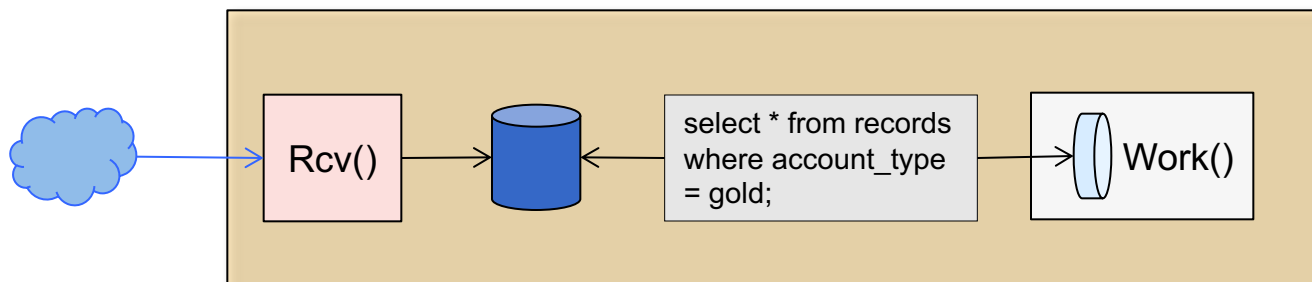
- Queuing of requests provides
  - ◆ Preserves client's asynchronous ability without losing requests
  - ◆ Receiving (Server) can process messages to its ability; rate (i.e. messages/sec)



Should **this** be synchronous or asynchronous?

# Enhancement 1: dequeue query to provide selective processing

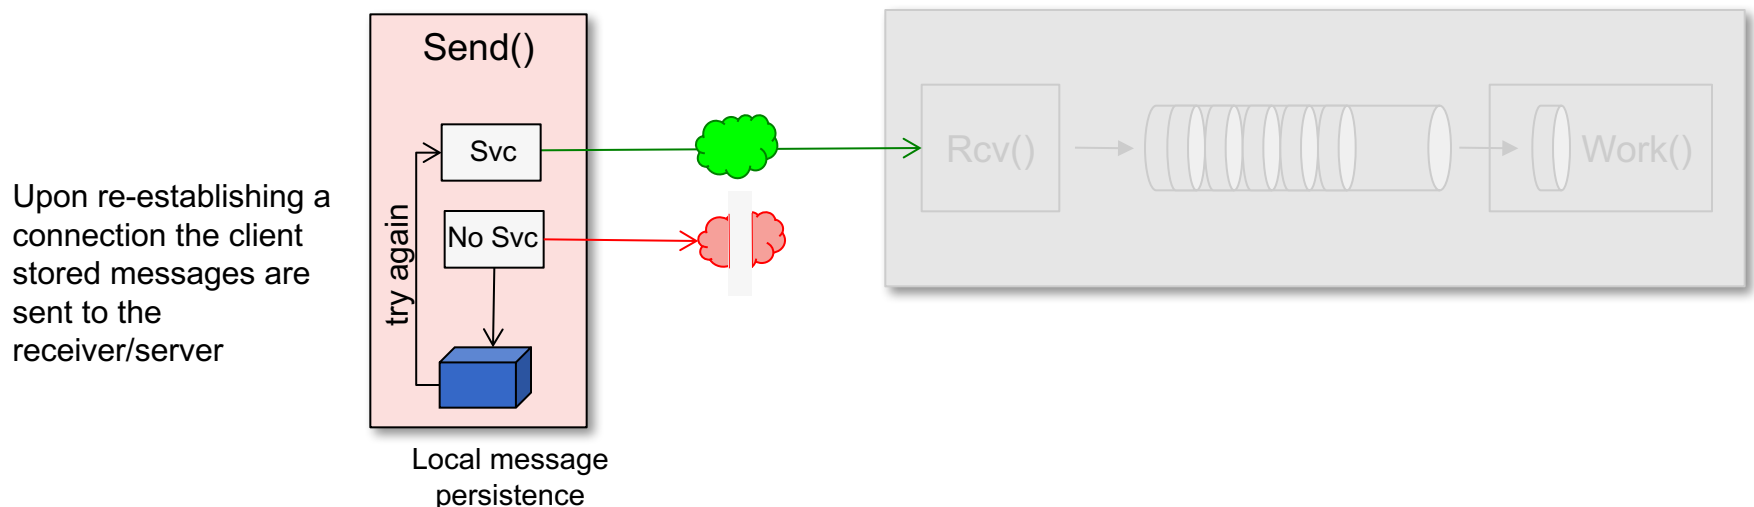
- enqueued requests in a basic queue acts as a FIFO
  - ♦ PRO: request processed as received, this will support passive sequencing (no special processing for sequentially dependent data)
  - ♦ CON: No (built-in) selective processing
    - E.g., prioritization of messages
- Query queue for pre-processing prioritization
  - ♦ Implementations can range from volatile memory to persistent message storage (use of SQL)



Incidentally, how do you imagine JMS works?

# Enhancement 2: client-side or middleware queuing (store-and-forward)


- Store-and-forward is the ability for the client-side (sender) to store messages if the receiver (server) is not available
  - ♦ From the client's perspective the message was sent (no action required under the assumption of asynchronous behavior)
  - ♦ Questions:
    - How would a synchronous communication fair?
    - What are the risks of this approach



# Enhancement 3: Quality of Service (QoS)


- What do we mean by QoS?
- What is our QoS design?

# QoS: Simple queuing




- Simple queue
  - ♦ “Best effort” or “Store-and-Forward”
  - ♦ Provides buffering between producer and consumer (overflow)
  - ♦ Simple to step up, use, and support under the right conditions
  - ♦ FIFO behavior
  - ♦ Limited to no control of resources – reactive or push design
    - lacks prioritization of requests
    - Prolific producers can starve other producers (clients)
  - ♦ Variants
    - Durable – storage provides buffer overflow and persisting for failure/recovery
    - Consumer pool – multiple consumers (I.e., MDB)

# QoS: Fair Queuing



- Fair Queuing (FQ) and variants
  - ♦ Another “Best effort” QoS
  - ♦ Use of multiple queues to bin requests
  - ♦ Simple spraying or overflow binning across multiple queues
- Variants
  - ♦ Employs simple algorithm(s) to determine which queue (bucket) to place a message
  - ♦ Weighted Fair Queuing (WFQ)
  - ♦ Hierarchical Weighted Fair Queuing (HWFQ)
  - ♦ Class-based WFQ (CBWFQ)

# QoS: Early detection (rejection)




- Random Early Detection (RED)
- Variants
  - ♦ Weighted RED (WRED)
    - Moves buffering and rescheduling onto the producer (client)
    - Drop messages when unable to process
    - Weighted – apply criteria (producer, priority, etc) to drop
    - Generally, not a friendly approach
  - ♦ Class-based WFQ (CBWFQ)
  - ♦ RSVP or Guaranteed service (hard QoS)
    - Reservation of queue



# Enhancement 4: Failover and scaling

- What is happening here?
- PROs and CONs?



# Other enhancements

- Concurrent/Parallel message processing
- Quality Control (pre-filtering)
- Dynamic payload sizes (optimizing throughput)
- Sequencing (Ordering) of messages
- Authorization and bandwidth management
  - ♦ QoS application
- Auditing
- Encryption and compression

# Many F/OSS and COTS packages

- Implementations

- ◆ OpenJMS (<http://openjms.sourceforge.net/>) – low activity since 2007
- ◆ NATS cloud mq – <https://nats.io>
- ◆ SwiftMQ (<http://www.swiftmq.com>)
- ◆ RabbitMQ, an AMPQ implementation
- ◆ ZeroMQ (<http://zeromq.org>)
- ◆ Apache
  - Qpid – AMPQ implementation
  - ActiveMQ (<http://activemq.apache.org/>)
- ◆ Tibco broker (<http://www.tibco.com>)

# Reading and references

- Papers and discussions
  - ♦ <http://bravenewgeek.com/dissecting-message-queues/>
  - ♦ <http://www.dynamicobjects.com/papers/w4spot.pdf>
  - ♦ <https://www.sdn.sap.com/irj/sdn/go/portal/prtroot/docs/library/uuid/50b7ac8d-0aed-2a10-d290-b64f44c4c1a9>
  - ♦ <http://www.precisejava.com/javaperf/j2ee/JMS.htm>
- Software used in to support the messaging lab (distributed off-line)
  - ♦ RabbitMQ – <http://www.rabbitmq.com>
  - ♦ Erlang – <http://www.erlang.org>

# Reading and references

- JMS

- ♦ <http://72.5.124.55/j2ee/1.4/docs/tutorial-update6/doc/JMS3.html>
- ♦ JMS Specification, version 1.1 available from <http://java.sun.com/products/jms/docs.html>

- Consensus

- ♦ Distributed Algorithms, Nancy Lynch, 1996
- ♦ Using Paxos to Build a Scalable, Consistent, and Highly Available Datastore, Rao, Shekita, Tata, 2011
- ♦ A Survey of Consensus Problems in Multi-agent Coordination, Ren, Beard, Atkins, 2005
- ♦ The Byzantine Generals Problem, Lamport, Shostak, Pease, 1982
- ♦ Paxos Made Simple, Lamport, 2001
- ♦ Paxos Made Live – An Engineering Perspective, Chandra, Griesemer, Redstone, 2007

## Appendix: JMS Summary

# Java Messaging Service (JMS)

- Java Messaging Service (JMS)
  - ♦ <http://java.sun.com/products/jms/index.jsp>
  - ♦ Specification is surprisingly readable
- JMS version 1.1 (2002)
  - ♦ J2EE 1.4 (2003)
  - ♦ Compatible with JMS 1.02
  - ♦ Significant restructuring to simplify API
  - ♦ Domain unification
    - Simplified API for general, reusable code
    - Common interface allows the same class to send and receive messages
    - Unification allows a session within a single TX to send and receive
    - Session can manage queues and topics

# Asynchronous-ish messaging


- JMS defines asynchronous messaging as independent request and response between the producer and consumer. The consumer can choose between blocking and event processing.
- Producer (sender)
  - ♦ Non-blocking – `MessageProducer.send()`
- Consumer (receiver)
  - ♦ Blocking
    - `MessageConsumer.receive()`
    - Consumer blocks on the `receive()` waiting for messages or time out
  - ♦ Non-blocking
    - `MessageConsumer.setMessageListener(MessageListener l)`
    - Implementation of `javax.jms.MessageListener` interface is invoked upon receipt of a `javax.jms.Message`
    - This is how MDBs are defined in Java EE

```
public class MyListener implements MessageListener {  
    public void onMessage(Message message) {...}  
}
```



# JMS Queues


- Point-to-Point (PTP)
  - ♦ Each message is delivered to one consumer
  - ♦ If no consumers are registered the message is held unless an expiration was set – `MessageProducer.setTimeToLive(msec)`
- `javax.jmx.QueueBrowser`
  - ♦ Look at messages without removing them
- Sequence
  - ♦ Producer connects to a message queue
  - ♦ Producer sends message (blocking)
  - ♦ Returned successful – message delivered to JMS server
  - ♦ If no consumers are listening, the message is held until a consumer connects
  - ♦ If more than one consumer is registered, the message is only delivered to one consumer



# JMS Topics


- Publish/Subscribe messaging (Fanout or Broadcast)
  - ◆ Each message can have multiple consumers (same message will not be delivered twice\*)
  - ◆ Messages are guaranteed to be delivered to any consumers if a consumer's subscription is durable. Otherwise, they are discarded.

```
Session.createDurableSubscriber()
```



*\*Unless transactional control is implemented and a rollback is performed*

# JMS supports the transport of many payload representations



- TextMessage
  - ♦ Text (a.k.a String)
- ObjectMessage
  - ♦ Serializable object
- StreamMessage
  - ♦ Sequence (FIFO) of java primitives
- BytesMessage
  - ♦ Stream of uninterpreted bytes
- MapMessage
  - ♦ Send name–value pairs
  - ♦ Names are String objects and values are primitives

# Example: Queue producer

```
// OpenJMS
Context context = null;
Connection connection = null;
try {
    context = new InitialContext();
    ConnectionFactory factory =
        (ConnectionFactory) context.lookup(factoryName);
    Destination dest = (Destination) context.lookup(destName);
    connection = factory.createConnection();

    Session session =
        connection.createSession(false, Session.AUTO_ACKNOWLEDGE);

    MessageProducer sender = session.createProducer(dest);

    connection.start();

    TextMessage message = session.createTextMessage();
    message.setText("Hello");
    sender.send(message);

} catch (Exception exception) {
    .
    .
    .
} finally {
    context.close();
    connection.close();
}
```

# Example: Queue consumer

```
Context context = null;
Connection connection = null;
try {
    context = new InitialContext();

    ConnectionFactory factory = (ConnectionFactory)
        context.lookup(factoryName);
    Destination dest = (Destination) context.lookup(destination);
    connection = factory.createConnection();

    Session session = connection.createSession(false,
        Session.AUTO_ACKNOWLEDGE);

    MessageConsumer receiver = session.createConsumer(dest);

    connection.start();
```

```
    Message message = receiver.receive();
    if (message instanceof TextMessage)
        TextMessage text = (TextMessage) message;
```

Should be contained in a  
loop: while (true) {...}

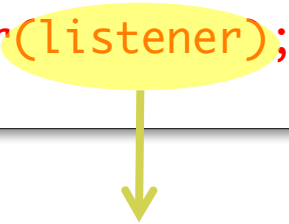
```
} catch (Exception exception) {
    exception.printStackTrace();
} finally {
    context.close();
    connection.close();
}
```

# A consumer that is event driven

```
// create the connection
connection = factory.createConnection();
session = connection.createSession(false,
    Session.AUTO_ACKNOWLEDGE);

MessageConsumer receiver = session.createConsumer(dest);

// register a listener
receiver.setMessageListener(listener);
```



```
public class MyListener implements MessageListener {
    public void onMessage(Message message) {
        if (message instanceof TextMessage) {...}
    }
}
```

# Message delivery is not guaranteed (perspectives)

- Message delivery mode
  - ◆ JMS supports two delivery modes
    - **NON\_PERSISTENT** – Does not require the JMS server to ensure message delivery
      - Delivery is **at-most-once** (which includes zero)
    - **PERSISTENT** – JMS server (provider) is instructed to ensure message is not lost between producer and consumer
      - Delivery is guaranteed **once-and-only-once** (successful/acknowledged)
      - However, the JMS server may “experience resource limitations” that could result in lost messages
- Message’s Time-To-Live
  - ◆ A producer can specify how long a message should be retained before discarding
  - ◆ A discarded message is not delivered (persistent or non-persistent)

# Acknowledgement modes when sending messages

- The three acknowledgement modes are
  - ◆ **AUTO\_ACKNOWLEDGE**
    - Automatically handled by the Session when
      - `send()` or `receive()` returns
      - Call to the `MessageListener.onMessage()` is called
  - ◆ **DUPS\_OK\_ACKNOWLEDGE**
    - Lazy session acknowledgement
    - Duplicate acknowledgements may be sent
  - ◆ **CLIENT\_ACKNOWLEDGE**
    - Explicit control of acknowledgement, consumer controlled – allows for batch processing and rollback



# The JMS Message class

- `javax.jms.Message`
  - ◆ Composed of
    - Header – methods beginning with `setJMS` and `getJMS`
    - Properties – `get/setXXXProperty()`
    - Body – implementation specific (e.g., `setText()`, `getText()`)

Sound familiar?

# Transacted sessions

- Transacted sessions are used with `CLIENT_ACKNOWLEDGEMENT` sessions to confirm or reject message handling
  - ◆ `Session.commit()`
  - ◆ `Session.rollback()`
- How it works
  - ◆ For a producer, messages are not sent to a consumer until `commit()` is called
  - ◆ For a consumer messages are not confirmed received until `commit()` is called

# Advantages and disadvantages in simple queuing

- Pros
- Cons

